Looking for Lacunae in Bitcoin Core's Fuzzing Efforts

Alex Groce
Northern Arizona University
United States

ABSTRACT

Bitcoin is one of the most prominent distributed software systems in the world, and a key part of a potentially revolutionary new form of financial tool, cryptocurrency. At heart, Bitcoin exists as a set of nodes running an implementation of the Bitcoin protocol. This paper describes an effort to investigate and enhance the effectiveness of the Bitcoin Core implementation fuzzing effort. The effort initially began as a query about how to escape *saturation* in the fuzzing effort, but developed into a more general exploration once it was determined that saturation was largely illusory, a byproduct of the (then) fuzzing configuration. This paper reports the process and outcomes of the two-week focused effort that emerged from that initial contact between Chaincode Labs and academic researchers. That effort found no smoking guns indicating major test/fuzz weaknesses. However, it produced a large number of additional fuzz corpus entries to add to the Bitcoin QA assets, clarified some long-standing problems in OSS-Fuzz triage, increased the set of documented fuzzers used in Bitcoin Core testing, and ran the first (to our knowledge) mutation analysis of Bitcoin Core's fuzz targets, revealing opportunities for further improvement. We contrast the Bitcoin Core transaction verification testing with that for other popular cryptocurrencies. This paper provides an overview of the challenges involved in improving testing infrastructure, processes, and documentation for a highly visible open source target system, from both the state-of-the-art research perspective and the practical engineering perspective. One major conclusion is that for well-designed fuzzing efforts, improvements to the *oracle* side of testing, increasing invariant checks and assertions, may be the best route to getting more out of fuzzing.

CCS CONCEPTS

• **Software and its engineering** → **Dynamic analysis**; **Software testing and debugging**.

KEYWORDS

fuzzing, saturation, test diversity, mutation analysis

1 INTRODUCTION

Bitcoin [26] is the most popular cryptocurrency, and one of the most visible (if controversial) "new" systems based on software to rise to prominence in the last decade. As we write, while volatile, Bitcoin consistently has a market cap of over half a trillion dollars since January of 2021. Due to its distributed, decentralized nature, Bitcoin in some sense is the sum of the operations of the code executed by many independent Bitcoin nodes, especially nodes that mine cryptocurrency. Bitcoin Core (https://github.com/Bitcoin/Bitcoin) is by far the most popular implementation, and serves as a reference for all other implementations. To a significant degreee, the code of Bitcoin Core *is* Bitcoin. The main Bitcoin Core repo on GitHub has over 57,000 stars, and has been forked more than 30,000 times.

Because of its fame and the high monetary value of Bitcoins, the Bitcoin protocol and its implementations are a high-value target for hackers (or even nation states interested in controlling cryptocurrency developments). Therefore, testing the code is of paramount importance, including extensive functional tests and aggressive *fuzzing*. This paper describes a focused effort to identify weaknesses in, and improve, the fuzzing of Bitcoin Core.

2 INITIAL CONTACT AND THE PROBLEM OF SATURATION

Chaincode Labs (https://chaincode.com/) is a private R&D center based in Manhattan that exists solely to support and develop Bitcoin. In March of 2021, Adam Jonas, the head of special projects at Chaincode, contacted the first author to discuss determining a strategy to improve the fuzzing of Bitcoin Core. In particular, at the time, it seemed that the fuzzing was "stuck": neither code coverage nor found bugs were increasing with additional fuzzing time. After some discussion, an 80 hour effort was determined as a reasonable scope for an external, research-oriented, look at the fuzzing effort. Before that effort, conducted over the summer of 2021, began, the problem of saturation resolved itself. Nonetheless, the issue that drove the initial desire for a researcher investigation is well worth examining. Moreover, understanding why Bitcoin Core fuzzing was, temporarily but not fundamentally, saturated, may be useful to help other fuzzing campaigns avoid the same false saturation problem.

Saturation, as defined in the blog post (https://blog.regehr.org/archives/1796) that brought Chaincode Labs to the first author, is when "We apply a fuzzer to some non-trivial system... initially it finds a lot of bugs... [but] the number of new bugs found by the fuzzer drops off, eventually approaching zero." That is, at first a particular fuzzer[1] applied to a system will tend to continuously,

---

[1]The term "fuzzer" can apply to either a fuzz engine such as libFuzzer or to a "fuzz harness/target" for the fuzzed program [15]. By "fuzzer" we mean a particular *fuzz engine*, though the definition of saturation works in either case.

sometimes impressively, increase both coverage and discovery of previously-unknown bugs. But, at some point, these bugs are known (and often fixed) and the fuzzer stops producing new bugs. Code and behavioral coverage seems to be *saturated*.

The underlying reason for saturation is that any fuzzer (or other test generator) explores a space of generated tests according to some, perhaps very complex, probability distribution. Some bugs lie in the high-probability portion of this space, and other bugs lie in very low probability or even zero probability parts of the space. Unsurprisingly, eventually the high probability space is well-explored, and the remaining bugs are found only very infrequently, if ever. The underlying empirical facts are cruel, as noted by Böhme and Falk: "[W]ith twice the machines, we can find *all known* bugs in half the time. Yet, finding linearly *more* bugs in the same time requires exponentially more machines" [3].

Chaincode Labs saw that code coverage and bugs were not increasing in their fuzzer runs, and wanted to break out of the saturation trap. The blog post suggested several methods for doing just that, and this paper explores some of them. Note that even though the fuzzing was not saturated, the methods for escaping saturation are also useful things to try in any fuzzing effort where finding as many bugs as possible is actually important.

2.1 We Hold the World Ransom For... ONE MILLION FUZZER ITERATIONS

One problem for inexperienced fuzzer users is that it is not always clear just how long a fuzzer needs to run. Anyone used to popular random testing tools, such as QuickCheck [7], may expect a "reasonable" budget to be on the order of hundreds of tests or at most a few minutes [21]. The fuzz testing research community has settled, to some extent, on 24 hours as a basis for evaluation of fuzzers [23], but even this may be considered a low-budget run in a serious campaign! The Solidity compiler fuzzing effort discussed in the saturation blog post (https://blog.trailofbits.com/2020/06/05/breaking-the-solidity-compiler-with-a-fuzzer/, https://blog.trailofbits.com/2021/03/23/a-year-in-the-life-of-a-compiler-fuzzing-campaign/) found many bugs only after running a specialized version of AFL for over *a month*.

One thing that quickly emerged from discussions with Chaincode before the primary 80 hour effort was the limited extent of the fuzzer runs performed in early April. The fuzzing includes a large number of targets, each with its own fuzz harness and executable. At the time, the basic strategy was to run libFuzzer on each of these for 100,000 iterations. Because some targets are very fast and a few, such as full message processing, are slow, this meant in practice fuzzing most targets for only 30-90 seconds, and even the slowest targets for only a little over an hour. The total time for over 100 targets was not negligible, but expecting such short runs for each target, after an initial exploration of the easy part of the probability space, to gain coverage or bugs very often, was simply unrealistic. In particular, for complex, critical targets such as transaction verification and end-to-end message processing, 100,000 iterations was as inadequate as Dr. Evil's famous ransom demand for "one million dollars" in the first Austen Powers movie. Neither the Chaincode team nor Dr. Evil was being unreasonable; they were both simply unaware of the "inflation" of needed resources required for serious fuzzing or blackmail of the world.

The first suggestion for escaping coverage, therefore was very simple: run the fuzzer longer! The Chaincode team immediately tried increasing their configuration to 5 million iterations, multiplying the number of executions, and runtime, by a factor of 50. Based on initial success with a few targets, this was done for all targets, and eventually became the new default. To a large extent, saturation was no longer a problem. This exploration of simply increasinig the fuzzing budget came early, about 15 days after the initial contact. The Chaincode team also added new seeds by, as advised, running more fuzzers, including AFL and Honggfuzz, in the same time frame.

By May 20th, Bitcoin Core was also in OSS-Fuzz (it was not at the time of the first discussions, due to reporting requirements, but negotiations settled this problem): https://github.com/google/oss-fuzz/tree/master/projects/bitcoin-core. From then on, Bitcoin Core has essentially been continuously fuzzed, and OSS-Fuzz quickly produced new crashes to investigate, and continues to do so: https://bugs.chromium.org/p/oss-fuzz/issues/list?q=bitcoin.

3 ADDING FUZZER DIVERSITY: USING ECLIPSER AND TRYING ENSEMBLE FUZZING

The most obvious solution when fuzzer A (e.g., libFuzzer) faces saturation on a target codebase is to bring in fuzzer B (e.g., Honggfuzz). In fact, the original blog post that attracted Chaincode's attention mentions this in the very definition of saturation: "Subsequently, a different fuzzer, applied to the same system, finds a lot of bugs." Of course, this is the easy case, and there can also be saturation across the meta-fuzzer defined as a composition of all fuzzers applied! Still, "throw in another fuzzer" is almost always a good idea when performing serious fuzzing campaigns. If that fuzzer adds some elements that are not present in any of the previously applied fuzzers, so much the bettter.

3.1 Fuzzing with Eclipser

Eclipser [6] is a fuzzer that combines AFL-like[2] coverage-driven mutation-based fuzzing with a scalable form of grey-box concolic testing. AFL++ [10] and libFuzzer https://llvm.org/docs/LibFuzzer.html, the two fuzzers most aggressively applied to the Bitcoin Core code, do not perform any kind of symbolic or concolic testing. Eclipser is fairly easy to apply to new systems, since it uses QEMU to run on an ordinary binary, so the first order of business was to see if it could find new bugs or at least cover code AFL++ and libFuzzer were not able to reach.

Either directly (from runs using Eclipser itself) or indirectly (performing long AFL or libFuzzer runs seeded with Eclipser-generated tests), Eclipser produced thousands of new corpus seed files that were accepted as PRs to merge into the Bitcoin Core QA assets repo https://github.com/bitcoin-core/qa-assets. The first author is now the third-largest contributor to the QA assets repo, in fact. These files added hundreds of new coverage edges to the basis for OSS-Fuzz and other testing of Bitcoin Core.

---

[2]In fact, the latest version of Eclipser simply uses AFL for non concolic fuzzing.

On the other hand, these files added *no* new covered statements; the additions were all edges, not exploration of completely uncovered code, the most promising form of new coverage [1, 13]. Nor did these tests expose any new bugs. Of course, in the long run testing in OSS-Fuzz or on the QA team's servers may use these tests to help cover new code or find new bugs, but the immediate impact is more modest.

One of the outcomes of the 80 hour effort was full documentation of how to run Eclipser v1.x on the Bitcoin Core code: https://github.com/bitcoin/bitcoin/blob/master/doc/fuzzing.md#fuzzing-bitcoin-core-using-eclipser-v1x.

3.2 Ensemble Fuzzing

Given that Eclipser, while useful, did not produce dramatic results, simply running and adding instructions for more fuzzers did not seem like a particularly productive use of much more of the two-week effort. However, because Eclipser was somewhat useful, one obvious way to exploit multiple fuzzers was to try *ensemble fuzzing*.

Ensemble fuzzing [5] is an approach that recognizes the need for diverse methods for test generation, at least in the context of fuzzing; using multiple fuzzers to seed each other and avoid saturation is a core motivation for ensemble fuzzing. Inspired by ensemble methods in machine learning [8], ensemble fuzzing runs multiple fuzzers, and uses inputs generated by each fuzzer to seed the other fuzzers. Ensemble fuzzing is currently (in principle) supported by the Enfuzz website (http://wingtecher.com/Enfuzz) and by the DeepState [12] front-end[3].

However, the EnFuzz website has never, since the Bitcoin Core effort began, actually been up and working, and various build and library version errors made it impossible to get the GitHub version (https://github.com/enfuzz/enfuzz) to work, either. DeepState's ensemble code is more limited, but works. However, it would require re-writing Bitcoin Core test harnesses to use DeepState's GoogleTest-like API. In theory, this is not a huge burden: the custom fuzzing core API built by the Bitcoin Core team is fairly similar, in concept, to DeepState, providing ways to obtain values of various types using different fuzzer back-ends. However, there are numerous small differences and specialized behaviors in the Bitcoin code that would require substantial work to re-engineer, and DeepState would add a significant build dependency to the Bitcoin testing infrastructure; such a change might well not make it through the Bitcoin Core approval process, once done. Further, the effort to re-write more than 100 targets would be substantial, even if most targets were fairly trivial to translate. Building a one-off version not intended to be rolled into the main codebase might be worthwhile, but would clearly take more than the 80 hours available, for unknown payoff. Therefore performing ensemble fuzzing of Bitcoin Core remains future work for the QA team, or awaits the arrival of more reliable ensemble fuzzers that operate through a libFuzzer, AFL, or Honggfuzz interface that requires no changes on the Bitcoin Core side.

Note that in one important sense Bitcoin Core *is* using ensemble fuzzing, in that OSS-Fuzz runs multiple fuzzers, including libFuzzer, AFL, and honggfuzz) with different compilation flags and sanitizers.

[3]See https://blog.trailofbits.com/2019/09/03/deepstate-now-supports-ensemble-fuzzing/.

```
namespace NetMsgType {
  extern const char* VERSION;
  extern const char* VERACK;
  extern const char* ADDR;
  extern const char *ADDRV2;
  extern const char *SENDADDRV2;
  extern const char* INV;
  extern const char* GETDATA;
  extern const char* MERKLEBLOCK;
  extern const char* GETBLOCKS;
  extern const char* GETHEADERS;
  extern const char* TX;
  extern const char* HEADERS;
  extern const char* BLOCK;
  extern const char* GETADDR;
  extern const char* MEMPOOL;
  extern const char* PING;
  extern const char* PONG;
  extern const char* NOTFOUND;
  extern const char* FILTERLOAD;
  extern const char* FILTERADD;
  extern const char* FILTERCLEAR;
  extern const char* SENDHEADERS;
  extern const char* FEEFILTER;
  extern const char* SENDCMPCT;
  extern const char* CMPCTBLOCK;
  extern const char* GETBLOCKTXN;
  extern const char* BLOCKTXN;
  extern const char* GETCFILTERS;
  extern const char* CFILTER;
  extern const char* GETCFHEADERS;
  extern const char* CFHEADERS;
  extern const char* GETCFCHECKPT;
  extern const char* CFCHECKPT;
  extern const char* WTXIDRELAY;
}
```

Figure 1: The message types declaration in `src/protocol.h`.

Additionally, the Bitcoin Core team has servers running different fuzzers. All of these are coordinated via the `qa-assets` repository to which our Eclipser-based tests were added. This is, however, a more manual and less controlled process than true ensemble cross-seeding on-the-fly during a fuzzing campaign, and there are suggestions that a well-chosen coordination strategy can significantly improve ensemble effectiveness.

4 TRYING SWARM FUZZING

Swarm testing [20] is a method for improving test generation that relies on identifying *features* [19] of tests, and disabling some of the features in each test. For instance, if features are API calls, and we are testing a stack with push, pop, top, and clear calls, a non-swarm random test of any significant length will contain multiple calls to *all* of the functions. In swarm testing, however, for each test *some of the calls (with probability usually equal to 0.5 for each call) will be disabled*, but *different calls* will be disabled for *each generated test*. This produces less variance between calls *within tests*, but much more variance *between tests*. Practically, in the stack example, it will enable the size of the stack to grow much larger than it ever would have any chance of doing in non-swarm testing, due to some tests omitting pop and/or clear calls. Swarm testing is widely used in compiler testing [24] and is a core element of the testing for FoundationDB [28].

Some Bitcoin Core fuzz harnesses resemble API-sequence generation. In particular, the process_messages harness produces method sequences of a fixed length, where each message is of a particular type (there are 34 distinct message types, again similar to an API-call "fuzz surface"; see Figure 1). Swarm testing is an obvious candidate approach for enhancing this important fuzz harness.

Inspection of the `process_messages` code reveals that not only does the harness (https://github.com/bitcoin/bitcoin/blob/master/src/test/fuzz/process_messages.cpp) not perform swarm testing, it doesn't even pick from the message types; it simply generates an arbitrary string that will be parsed as a message type, a fuzzer-controlled arbitrary number of times:

```
while (fuzzed_data_provider.ConsumeBool()) {
  const std::string random_message_type{
    fuzzed_data_provider.ConsumeBytesAsString(
      CMessageHeader::COMMAND_SIZE).c_str()};
```

DeepState [12] (https://github.com/trailofbits/deepstate) provides strong support for swarm fuzzing. In DeepState, transforming the call to `ConsumeBytesAsString` to a `OneOf` nondeterministic choice operator over the strings defined in Figure 1, and compiling the harness with `-DDEEPSTATE_PURE_SWARM` would enable automatic fuzzer controlled swarm testing. However, as with ensemble fuzzing, while there are similarities to the Bitcoin Core infrastucture for fuzz targets and the DeepState API, re-writing the fuzz targets to use DeepState seemed, and change many string generation calls to `OneOf` selections was too large a time investment for the likely payoff without first trying a less ambitious approach.

The solution was to focus on developing a swarm harness for the most promising target, `process_messages`, alone, and determine if an aggressive campaign on that target would produce substantial new coverage, or even unknown bugs. This `process_messages_swarm` harness required modest effort, and could show if the strategy was worth devoting more resources to applying more broadly. After four weeks of fuzzing, however, the new coverage generated was minimal; less than would be expected from simply devoting that level of effort to the original `process_messages` target. Why?

First, it is likely that the automatic translation of more than 10,000 corpus files for `process_messages` was imperfect. Parsing raw bitstreams to find commands is hard, and writing a dynamic parser to use in-flight data to extract the meaning of each test was not feasible in the 80 hour effort. This probably meant some corpus files lost some or all of their value, leaving the new fuzz effort behind the equivalent `process_messages` fuzz.

Most importantly, however, the `process_messages` corpus is regularly *cross-seeded* with data from specialized fuzz targets that generate only one type of message, e.g., such as the tests in https://github.com/bitcoin-core/qa-assets/tree/main/fuzz_seed_corpus/process_message_mempool, which correspond to `extern const char* MEMPOOL` in Figure 1 . This not only provides some of the power of swarm testing, it achieves a second goal of the swarm harness. Namely, one concern was that since `process_message` and `process_messages` generate raw strings that may not even have a valid message type, the fuzzers spend too much effort fuzzing the type to little benefit. The specialized targets avoid this problem by constraining the fuzzing to only generate one type of message when running in the specialized mode. The ability to run `process_message` with a single message type and the frequent introduction of the resulting inputs into `process_messages` (and the generic, unconstrained `process_message` fuzzing) probably, in a less automated way, also achieves many of the benefits of swarm testing itself: mixing complex lengthy runs of a single type or mix of types. Where introducing "real" swarm testing is difficult, but fuzzing infrastructure supports this mode of operation, it may be a useful alternative. It is on the other hand unclear how many fuzz frameworks are as sophisticated as Bitcoin Core's, making this possible; in many cases simply implementing swarm directly (or writing the harnesses using DeepState) would likely be easier. It does show that some advanced fuzzing strategies can be anticipated by particularly savvy and capable fuzz engineers, willing to directly use raw fuzzing data, and write tools to support that kind of low-level hand-tuning of fuzz corpuses.

5 SIDE ISSUES: SPURIOUS BUGS, FUZZER MYSTERIES, AND AFL STABILITY

A critical point about testing is that it is like other kinds of software development. In practice, while adding "features" or at least optimizations of existing features — in our context, adding new fuzzers or cutting-edge methods such as swarm testing — is usually seen as the most interesting aspect of the work, the reality is that developer/tester time is often spent on more mundane, frustrating issues.

For example, in the Bitcoin Core effort, early attempts to apply Eclipser and AFL ran into a problem: at some point both fuzzers would essentially stop working and begin producing a vast number of spurious crashes. By spurious crashes, we mean unreproducible crashes: inputs that the fuzzer marks as causing a crash, but which, when executed in a non-fuzzing context using the fuzz harness, do not crash. Spurious crashes are not unlike the pheonomenon of *flaky tests* [9], one of the banes of automated software testing. A flaky test is a test that, for the same snapshot of test code and code under test, sometimes fails and sometimes passes. In fact, arguably, spurious crashes in fuzzing are *flaky tests* where the flakiness is due to an environmental dependence: the test fails when run under the fuzzer. Flaky tests and spurious crashes consume signficant testing and developer/test engineer resources. During the very first discussions with the Chaincode team, when Chaincode first ran AFL, some spurious crashes were produced; these were eventually ignored. This happens to many spurious crashes, unlike flaky tests which manifest in standing regression tests that often have to be either removed or otherwise mitigated. In fuzzing, simpy throwing away spurious crashes without understanding them is a more feasible option. OSS-Fuzz runs on Bitcoin Core, using AFL, also produced some spurious crashes.

Unlike most spurious crashes, these were eventually (partly) understood. Because AFL and Eclipser, run more than a few hours, inevitably reached as stage where they *only* produced spurious crashes (instead of doing any useful fuzzing), the problem had to be understood, if these fuzzers were to be useful for serious fuzzing. The first author discovered that the AFL and Eclipser problems were, strictly speaking, not flaky or spurious at all. The fuzz harnesses all fail if there is insufficient storage space on `tmp` to execute the Bitcoin software. At some point, this happens, even if fuzzing begins with a large amount of free storage. The problem is that some (but not all) fuzz executions under AFL and Eclipser fail to clean up /tmp/test_common_Bitcoin Core; each run that fails to clean up leaves an 18-19 megabyte footprint. Over the millions of executions involved in a few hours of fuzzing, this inevitably fills the space.

This problem was also (probably) causing the original AFL issues seen by Chaincode, and (almost certainly, on examination of the OSS-Fuzz logs) the OSS-Fuzz failures that did not reproduce. Multiple possible solutions were discussed (at length) and proposed in a PR (https://github.com/bitcoin/bitcoin/pull/22472), but as we write none of these have been deemed successful. Also, the underlying reason why AFL and Eclipser sometimes fail to clean up /tmp is not understood; removing all threads did not fix the issue, though it might be useful for other reasons. The presence of threads may explain the relatively low stability of AFL fuzzing on Bitcoin Core, another issue raised during the 80 hour effort: https://github.com/bitcoin/bitcoin/issues/22551. Low stability indicates that when an input is executed multiple times, it does not always take the same path. The worse the stability, the less useful signal AFL is receiving from path coverage during fuzzing. Bitcoin Core's stability hovers around 80%, possibly as a result of the presence of threads in the code. Even if the threads do not cause nondeterminisic behavior of the Bitcoin Core system, from a functional point of view, they may cause path coverage to vary more than AFL "likes." It is hard to guess if this has a serious adverse effect on fuzzing performance.

One reason these issues have not been resolved is that they do not affect the most effective (thus far) fuzzing approach. Most new coverage obtained on Bitcoin Core, and most bugs found, can be credited to libFuzzer, which is not affected by the tmp problem, and at least fails to complain about stability, even if it is affected. At present, users of Eclipser on Bitcoin Core are advised to perhaps apply one of the patches proposed in the open PR on the problem, or to manually clean /tmp in some other way.

6 MUTATION ANALYSIS

Attempting to improve a fuzzing effort is one way to find problems with the effort; if you succeed, you found a weakness. However, none of the attempts exposed a serious problem. Adding more fuzzers would be *good*, but was not obviously *essential*. Ensemble fuzzing was not feasible, and swarm testing was, for practical purposes, already performed by alternative means. An alternative is to directly look for holes in testing. The Bitcoin Core fuzzing team clearly was measuring and inspecting code coverage (see https://marcofalke.github.io/btc_cov/), so little value would be added by inspecting traditional coverage alone. Mutation testing/analysis [27], however, subsumes code coverage and adds extremely valuable information on *oracle power* in addition to mere coverage [14]. This is perhaps especially valuable in fuzzing, where "you only see crashes" is a persistent concern.

We used the universal mutator (https://github.com/agroce/universalmutator) [17], a mutation tool already used widely in the blockchain and smart contract world, to mutate the Bitcoin and other popular cryptocurrencies core transaction verification code. We focused on transaction verification code, as it is generally well-covered by tests, and obviously an extremely critical functionality for the blockchain.

6.1 Mutation Testing Bitcoin

To perform mutation analysis on Bitcoin, we generated mutations for code in the tx_verify.cpp file. We then evaluated the ability

Figure 2: A comparison of mutation kills for `tx_verify.cpp` **when subjected to different testing and fuzzing parameters.**

of fuzzing and testing to detect such mutants. Figure 2 summarizes our mutation analysis of the file.

The universal mutator generated 430 compiling mutants of the file in less than three hours (out of caution, we forced a clean make of the full Bitcoin build for every mutant, which is expensive for compiling mutants). Two functions with interesting tests overlapped with our fuzzing targets in tx_verify.cpp code: process_message_tx and coins_view. Additional functions process_message and process_messages were also fuzz tested and of potential interest, but the relevant corpus entries were duplicated in process_message_tx (we verified these provided no additional mutant kills). We ran fuzzing for five minutes using lib-Fuzzer exploration based on the full (and large) QA asset corpus for each harness, with all sanitizers enabled. The process_message_tx target was able to detect 24 mutants, and the coins_view harness was able to detect 32 mutants, for a total of 50 mutants. In other words, slightly short of 12% over all the generated mutants. This is not a bad result: fuzzing inherently has trouble detecting subtle, non-crash-inducing, bugs in code, because writing a strong specification of correct behavior that covers all the bizarre and pointless inputs produced in fuzzing is often impractical, or would require a specification nearly as complex as the code itself. This is one reason *differential* fuzzing is promising: a reference implementation is such a specification. Bitcoin Core's cryptographic elements are, in fact, differentially fuzzed https://github.com/bitcoin/bitcoin/pull/22704#issuecomment-898989809.

A major purpose of fuzzing is, then, to address limits in more traditional functional testing, where known inputs are paired with expected behavior. While functional or unit testing is very powerful, the kinds of bugs found in vulnerabilities often involve the kind of inputs that don't appear in "normal" unit/functional tests, as shown by the success of fuzzing and security audits [16]. The real

question, then, is how many mutants that survive Bitcoin Core's extensive functional tests survive fuzzing.

The answer is: not too many. The functional tests without sanitizers enabled catch an additional 278 mutants. Turning on sanitizers (which is very expensive — we *only* ran it for mutants surviving all other tests, as indicated by the zero overlap and the asterisk in Figure 2) catches an additional 12 mutants. Only 90 of the 430 compiling mutants survive all tests, for an overall mutation score of 79.07%. On the other hand, fuzzing only adds two mutant kills to the functional testing.

Moreover, manual inspection of the surviving mutants showed that at least 29 of these were clearly semantically equivalent to the un-mutated code. For instance, many mutants removed or weakened an assertion; clearly this cannot ever be detected, since it can only transform failing tests into passing tests. The full, detailed list of surviving mutants, prioritized by an FPF ranking [4, 11], is available here: https://github.com/agroce/bitcorpus/blob/master/mutation/prioritized_full_inspect.txt. Discussion with the Bitcoin Core team is ongoing as we write (https://github.com/bitcoin/bitcoin/issues/22690), but thus far none of the surviving mutants seem to expose serious testing problems. After manual pruning, the mutation score is 85.8%, and some of the remaining 61 mutants are likely also equivalent. The limitations of the fuzzing oracle are clear: fuzzing covers 98% of statements and 72.6% of branches as we write, but can kill fewer than 12% of mutants. The much greater killing power of the functional tests obviously does not lie in the marginal 0.7 percentage points and 0.5 percentage points of statement and branch coverage obtained; it lies in the ability to reject incorrect executions that do not crash or set off a sanitizer alarm.

This raises the question: why fuzz? The coverage for high quality (if imperfect) functional tests such as those for Bitcoin Core will often be considerably higher, and the oracle will almost always be *much* more powerful. The answer lies in the fact that, even in the presence of such high quality tests, fuzzing uncovers subtle bugs that functional tests designed by humans will almost never detect, e.g. https://github.com/bitcoin/bitcoin/issues/22450[4]. Fuzzing is *not* a replacement for functional/unit tests; and functional/unit tests are not a replacement for fuzzing. In our mutation analysis, consider the two mutants detected by coins_view fuzzing alone. In the traditional, score-based, view of mutation analysis, the coins_view fuzz harness would be seen as performing badly. But it detects 2 (hypothetical) bugs not detectable by other means; in the real world, if one such bug is exploitable, detecting it may "pay for" all the fuzzing effort, and there will seldom be just one such bug (see https://github.com/bitcoin/bitcoin/issues?q=is%3Aissue+fuzz+is%3Aclosed+label%3ABug for an approximate list of fuzzer-detected, fixed bugs in Bitcoin Core). If the "good guys" don't fuzz well, you can be sure the bad guys will, for software protecting billions of dollars of assets.

```
validation.mutant.80.go: ../go-ethereum/signer/fourbyte/validation.go:79
*** Original
--- Mutant
***************
*** 76,82 ****
    switch {
    case tx.GasPrice == nil && tx.MaxFeePerGas == nil:
        messages.Crit("Neither 'gasPrice' nor 'maxFeePerGas' specified.")
!   case tx.GasPrice == nil && tx.MaxPriorityFeePerGas == nil:
        messages.Crit("Neither 'gasPrice' nor 'maxPriorityFeePerGas' specified.")
    case tx.GasPrice != nil && tx.MaxFeePerGas != nil:
        messages.Crit("Both 'gasPrice' and 'maxFeePerGas' specified.")
--- 76,82 ----
    switch {
    case tx.GasPrice == nil && tx.MaxFeePerGas == nil:
        messages.Crit("Neither 'gasPrice' nor 'maxFeePerGas' specified.")
!   /*case tx.GasPrice == nil && tx.MaxPriorityFeePerGas == nil:*/
        messages.Crit("Neither 'gasPrice' nor 'maxPriorityFeePerGas' specified.")
    case tx.GasPrice != nil && tx.MaxFeePerGas != nil:
        messages.Crit("Both 'gasPrice' and 'maxFeePerGas' specified.")
```

Figure 3: Mutation Example for Ethereum Transaction Validation

```
switch {
case tx.GasPrice == nil && tx.MaxFeePerGas == nil:
    messages.Crit("Neither 'gasPrice' nor 'maxFeePerGas' specified.")
case tx.GasPrice == nil && tx.MaxPriorityFeePerGas == nil:
    messages.Crit("Neither 'gasPrice' nor 'maxPriorityFeePerGas' specified.")
case tx.GasPrice != nil && tx.MaxFeePerGas != nil:
    messages.Crit("Both 'gasPrice' and 'maxFeePerGas' specified.")
case tx.GasPrice != nil && tx.MaxPriorityFeePerGas != nil:
    messages.Crit("Both 'gasPrice' and 'maxPriorityFeePerGas' specified.")
}
// Semantic fields validated, try to make heads or tails of the call data
db.ValidateCallData(selector, data, messages)
return messages, nil
```

Figure 4: Coverage Example for Ethereum Transaction Validation

6.2 Other Cryptocurrency Projects

Our Bitcoin work prompted mutation analysis for other popular cryptocurrencies. We generated mutants using the universal mutator for code in Ethereum, Solana, Dogecoin, Avalanche, Stellar, Algorand, and Cosmos, a selection from the top 30 cryptocurrencies[5] Our selection is influenced by the availability of code coverage, and indicative of the further opportunities for mutation testing in Bitcoin and popular altcoins, rather than a comprehensive survey. Because it requires significant effort to set up or otherwise understand the extent of fuzz testing in this heterogeneous selection of projects, our mutation analysis here only considers pure mutation testing against the project test suite (no fuzzing).

For the same reasons above, we targeted transaction validation logic and identified candidate files by searching for keywords like transaction, verify, sign, and validate. We manually inspected functions and test coverage for these functions (where applicable) to identify which files would be interesting targets to generate mutants for. Ultimately we settled on one to three files per project that are representative of some interesting and tested functionality (a choice that we readily acknowledge is by no means comprehensive or suggestive of a project's quality and testing as a whole).

Experimental Setup. rvt says: a TODO here that answers: - what is the range of number of mutants we generated across all projects (min and max) - how long did mutant generation take? - how many

---

[4]Comments on this bug, such as "Another win for fuzzing, oh wow." and "Fuzzer rulez!" show that the Bitcoin Core team has little doubt about the power of fuzzing.
[5]By market capitalization according to https://coinmarketcap.com.

Project	File path	LOC	Mutation score	File coverage	Project coverage	Inception date
bitcoin	src/consensus/tx_verify.cpp	210	78.6%	98.7%	84.2%	12/2010
go-ethereum	core/block_validator.go	129	70.1%	81.0%		
	signer/fourbyte/validation.go	127	49.5%	60.0%	58.8%	12/2013
	signer/core/signed_data.go	1,044	25.3%	69.3%		
solana	perf/src/sigverify.rs	1,246	????%	74.48%		
	core/src/validator.rs	2,016	-	73.29%	82.2%	02/2018
	core/src/tvu.rs	494	-	63.12%		
dogecoin	src/bitcoin-tx.cpp	847	58.7%	-	70.1%	12/2013
avalanchego	vms/platformvm/add_subnet_validator_tx.go	308	57.3%	81.0%	63.6%	03/2020
stellar	src/historywork/VerifyTxResultsWork.cpp	192	85.1%	-	-	11/2014
go-algorand	ledger/eval.go	1,551	99.8%	86.0%	52.2%	06/2019
cosmos-sdk	x/auth/ante/sigverify.go	510	73.1%	-	-	02/2016

Table 1: Code Coverage and Mutation Scores Across Popular Cryptocurrencies. <u>Mutation score</u> represents the proportion of mutants that were killed divided by the total number of mutants (higher is better). <u>File coverage</u> and <u>Project coverage</u> report statement level coverage for the file and entire project, respectively. <u>LOC</u> represents the lines of code of the chosen file. Absent entries indicated by – means we could not obtain coverage data, or otherwise curbed additional experiments. **rvt asks: what do we want to say about inception date here, or in the text? i presume there's something interesting to say here since we're including the data...**

mutation operators (or "rules") were used? - where they the same or different for each project?

Experimental Results. rvt says: a TODO to address in this subsection, or elsewhere above as examples of mutation testing: a handful of example mutants and/or mutation operators (either those that were unkilled by the Bitcoin runs, or the ones we did in our altcoin experiments) in the form of a readable table or something like that. i think would be really interesting to the reader. e.g., "what kinds of mutation operators did we decide on? which kinds of mutants survived?". even if it's standard stuff, it's probably not something the reader wants to guess about, and really is a "oh that's interesting" kind of thing to look at. we have the figures but it's difficult to read and these are really trying to make a different point (See discussion) so a table with mutation operators/rules or examples would be awesome

Table 1 summarizes our results, which lists the **Mutation score**, **File coverage**, and **Project coverage** over the selected file. **Mutation score** represents the proportion of mutants that were killed by the project's test suite divided by the total number of mutants (higher is better). File and project metrics represent statement coverage. statement level. In principle, we expect that higher file coverage (i.e., more tested code) to correlate with a higher mutation score.don't know if this needs a cite but that seems intuitive, sort of?

Out of our selection, Bitcoin's mutation score ranks high. Our experience working with Bitcoin Core developers show that they are proactive about code quality and testing, which is likely to influence (whether directly or indirectly) mutation scores.question for Alex: are we comfortable saying something along these lines?

just an attempt to reword what the previous sentence here tries to convey. otherwise, let's cut it At the same time, code around transaction verification and validation may naturally vary depending on context, code organization, and language as reflected by the differences in lines of code in our selection (and unknown function dependencies or testing in the rest of the codebase). Thus, the extent of our mutation testing per file ultimately serve as individual datapoints that are difficult to compare horizontally across projects. rvt says: changed things because basically like, cool, bitcoin is high, but like, we're mutating a 210 line file and it has lots of calls to lots of other code and the mutation score could be lower or higher if we decided to mutation test something somewhere else along the callchain that a tx_verify operation operates on. we really really can't say with confidence or even speculate (IMO) that this mutation score is super meaningful if we're tempted to compare it to other projects. Not even Doge fork because it's not the same file? If it were the same file I might feel differently.

Investigating files with lower coverage (e.g., those in Ethereum), we noticed that many of the generated mutants removed lines related to error checking (e.g., in switch case statements) that were apparently not covered by tests (cf. Figures 3 and 4).

Considerations for further study. Unit tests and corresponding coverage reports, while useful in our context of mutation testing, contribute only to one aspect of code quality and ignores other kinds of testing. Incorporating integration and differential tests with coverage would deliver deeper insights of under-tested code and overall code quality.

In a similar vein, we limited our mutation testing to a handful of files based on a convenience sample. A deeper mutation analysis

```
FUZZ_TARGET(parse_script)
{
    const std::string script_string(buffer.begin(), buffer.end());
    try
        (void)ParseScript(script_string);
    catch (const std::runtime_error&)

}
```

Figure 5: Fuzz target for script parsing.

with a higher allocation of resources (perhaps justifying resources comparable to fuzz testing) poses interesting possibilities that may allow more direct comparison and contrast among projects.

7 HOW GOOD IS THE BITCOIN CORE TESTING AND FUZZING?

In a sense this is a fundamentally hard question to answer: certainly, an 80 hour effort by one fuzzing researcher, who was generally familliar with blockchain and Bitcoin technology, but not at all familliar with the Bitcoin Core implementation, in terms of high level approach, testing infrastructure, or (of course) detailed codebase, cannot provide anything like a definitive answer to that question.

The previous section discusses our limited effort to place the Bitcoin Core test effort in the context of other cryptocurrencies, but the limits of that evaluation make it hard to draw strong conclusions. The complexities of the test efforts and the various ways the systems could be compared make this a weak basis for evaluating Bitcoin Core itself. Bitcoin Core has about 10,000 lines of fuzz harness code; Go-Ethereum has about 74KLOC of fuzz harness code; is the fuzzing thus likely to be 7 times better? We don't claim to know. Instead, we base an evaluation on the effort to improve Bitcoin Core fuzzing, itself.

The initial fuzzing effort, at the time of contact, was certainly operating in a suboptimal way, and hence *seemed* to be facing a difficult saturation problem. However, the simplest of industrial-fuzzing-savvy advice (run the fuzzers longer, add the system to OSS-Fuzz, try more fuzzers) was sufficient to make the fundamentally sound fuzzing basis more useful.

Additional efforts to improve the fuzzing added some value in terms of new corpus entries, and paved the way for better documented, and more systematic, application of a variety of fuzzers. Fundamentally, however, these improvements were relatively small compared to the overall scope of Bitcoin testing. This was substantially due to the limited timeframe, but not due to either a lack of expertise or serious attempts to improve fuzzing, given that timeframe, we believe.

The basic cause of the limited improvement was that, at present, Bitcoin Core has a well-designed and effective fuzzing infrastructure, supported by intelligent manual augmentation of the more automated aspects of the fuzzing.

7.1 The Bitcoin Core Fuzzing Infrastructure

The basic fuzzing infrastructure for Bitcoin resides in src/tests/fuzz, which consists of about 200 fuzz target implementations, ranging in size from a few lines to about 500 lines, plus common infrastructure, primarily found in FuzzedDataProvider.h. The fuzzing implementation is compact, but supports a common way to write fuzz targets that perform fuzz generation of C++ generic and Bitcoin-specific data structures. The generators are designed to allow different fuzzers to serve as the "back end" of the system, as in DeepState. This approach allows some targets to be very compact, e.g., see Figure 5.

In two weeks, it is impossible to examine every fuzz target in detail, much less understand how they interact with the targeted code, but most inspected targets are well-designed, and in cases where targets appear to be too generic, wasting time, as with process_message, there seems to be manual seeding that avoids the problem of forcing fuzzers to generate too many magic strings.

One sign that the fuzzing is of reasonable quality is that the gap in code coverage between fuzzing and manually constructed, extensive, functional tests is not very large. For both fuzzing and functional tests, some missing coverage and some unkilled mutants are almost certainly (and stated to be such by Bitcoin Core team members to us) redundant or defensive code to protect against "impossible" problems. This is unsurprising in code intended to be highly reliable; coverage of core file system functionality was only 89% during a year of extensive random testing of a file system developed for the Curiosity Mars Rover performed by NASA/JPL's Laboratory for Reliable Software [18]. The missing coverage was mostly defensive, redundant code, including fail-safe termination after (presumably impossible without hardware failure) assertion failures.

7.2 Improving Bitcoin Core Fuzzing

The lack of a huge coverage gap, and the existence of a huge mutant-detection gap, between the fuzzing and the functional tests suggests the most promising method for improving the Bitcoin Core fuzzing: manual, expert (in the codebase) effort to improve the oracles used by existing fuzz targets, or efforts to craft custom, more restricted, fuzz targets with stronger oracles when this is not feasible.

Essentially, the lack of a huge coverage gap suggests that failure to explore the input space is not the biggest problem for the Bitcoin Core fuzzing; there's significant room for improvement, especially with some targets, but overall the coverage, especially for critical code, is good. What the mutation testing reveals is that many fuzz executions that would be rejected by the standards of the functional tests do not induce a crash.

Building fuzz harnesses with complex correctness checks is hard, of course; the functional tests know exactly what inputs are being provided to APIs, and can check for expected behavior. Trying to inject this kind of check into fuzz harnesses is hard. When applicable, more generic, "mathematical" constraints such as are used in property-based testing [7] can help, but these are often hard to apply at the end-to-end level of a fuzz harness, without turning the fuzz harnesses into a spaghetti code mess of checks for various properties of the inputs. In the worst case, these end up simply reflecting developer/tester notions that already found their way into the code itself. Differential testing [25] against other implementations might also help here, but is probably best done at a higher end-to-end level than inside these fuzz targets.

There is no easy solution to this problem, but the most promising route is probably to focus on adding invariant assertions to the non-test code itself; these assertions can be executed by both functional and fuzz tests, and avoid the problem of duplicating analysis of inputs. At present, the Bitcoin Core code has about 1,800 `assert` statements, over a codebase of about 180KLOC of C and C++. The resulting ratio of about one assertion per 100 lines of code is not terrible, but is at the lower limit of what many consider to be an acceptable assertion ratio for critical code. Given that Bitcoin Core has more than 1,800 functions with `void` type (functions that cannot provide a value to check in a higher function, other than a side effect) alone, the code base obviously doesn't meet the NASA/JPL proposal of having an average of two assertions per function in a critical codebase [22]. There are only five assert statements in the `src/consensus` directory, which has about 500 lines of code, suggesting that this ratio is not notably higher in the most critical code.

The conclusion of the 80 hour effort therefore, is that the most effective way to improve fuzzing at present might be not to focus on covering the code and state space, but to focus on increasing the oracle power of *all* Bitcoin Core testing. Arguably, the greatest weakness of traditional code coverage is that it focuses too much attention on the "input side" of testing and too little on the oracle side [2], which is easier for even dedicated testing efforts to neglect in the pursuit of covering every branch and path.

8 CONCLUSIONS

REFERENCES

[1] Iftekhar Ahmed, Rahul Gopinath, Caius Brindescu, Alex Groce, and Carlos Jensen. 2016. Can Testedness Be Effectively Measured?. In *Proceedings of the 2016 24th ACM SIGSOFT International Symposium on Foundations of Software Engineering* (Seattle, WA, USA) *(FSE 2016)*. Association for Computing Machinery, New York, NY, USA, 547–558. https://doi.org/10.1145/2950290.2950324

[2] Earl T Barr, Mark Harman, Phil McMinn, Muzammil Shahbaz, and Shin Yoo. 2014. The oracle problem in software testing: A survey. *IEEE transactions on software engineering* 41, 5 (2014), 507–525.

[3] Marcel Böhme and Brandon Falk. 2020. Fuzzing: On the Exponential Cost of Vulnerability Discovery. In *Proceedings of the 28th ACM Joint Meeting on European Software Engineering Conference and Symposium on the Foundations of Software Engineering* (Virtual Event, USA) *(ESEC/FSE 2020)*. Association for Computing Machinery, New York, NY, USA, 713–724. https://doi.org/10.1145/3368089.3409729

[4] Yang Chen, Alex Groce, Chaoqiang Zhang, Weng-Keen Wong, Xiaoli Fern, Eric Eide, and John Regehr. 2013. Taming Compiler Fuzzers. In *Proceedings of the 34th ACM SIGPLAN Conference on Programming Language Design and Implementation* (Seattle, Washington, USA) *(PLDI '13)*. Association for Computing Machinery, New York, NY, USA, 197–208. https://doi.org/10.1145/2491956.2462173

[5] Yuanliang Chen, Yu Jiang, Fuchen Ma, Jie Liang, Mingzhe Wang, Chijin Zhou, Xun Jiao, and Zhuo Su. 2019. Enfuzz: Ensemble fuzzing with seed synchronization among diverse fuzzers. In *USENIX Security Symposium*. 1967–1983.

[6] Jaeseung Choi, Joonun Jang, Choongwoo Han, and Sang Kil Cha. 2019. Grey-box Concolic Testing on Binary Code. In *Proceedings of the International Conference on Software Engineering*. 736–747.

[7] Koen Claessen and John Hughes. 2000. QuickCheck: a lightweight tool for random testing of Haskell programs. In *International Conference on Functional Programming (ICFP)*. 268–279.

[8] Thomas G Dietterich et al. 2002. Ensemble learning. *The handbook of brain theory and neural networks* 2 (2002), 110–125.

[9] Moritz Eck, Fabio Palomba, Marco Castelluccio, and Alberto Bacchelli. 2019. Understanding Flaky Tests: The Developer's Perspective. In *Proceedings of the 2019 27th ACM Joint Meeting on European Software Engineering Conference and Symposium on the Foundations of Software Engineering* (Tallinn, Estonia) *(ESEC/FSE 2019)*. Association for Computing Machinery, New York, NY, USA, 830–840. https://doi.org/10.1145/3338906.3338945

[10] Andrea Fioraldi, Dominik Maier, Heiko Eißfeldt, and Marc Heuse. 2020. AFL++ : Combining Incremental Steps of Fuzzing Research. In *14th USENIX Workshop on Offensive Technologies (WOOT 20)*. USENIX Association. https://www.usenix.org/conference/woot20/presentation/fioraldi

[11] Teofilo F. Gonzalez. 1985. Clustering to Minimize the Maximum Intercluster Distance. *Theoretical Computer Science* 38 (1985), 293–306.

[12] Peter Goodman and Alex Groce. 2018. DeepState: Symbolic unit testing for C and C++. In *NDSS Workshop on Binary Analysis Research*.

[13] Rahul Gopinath, Carlos Jensen, and Alex Groce. 2014. Code Coverage for Suite Evaluation by Developers. In *Proceedings of the 36th International Conference on Software Engineering* (Hyderabad, India) *(ICSE 2014)*. Association for Computing Machinery, New York, NY, USA, 72–82. https://doi.org/10.1145/2568225.2568278

[14] Alex Groce, Mohammad Amin Alipour, and Rahul Gopinath. 2014. Coverage and Its Discontents. In *Proceedings of the 2014 ACM International Symposium on New Ideas, New Paradigms, and Reflections on Programming & Software* (Portland, Oregon, USA) *(Onward2014)*. Association for Computing Machinery, New York, NY, USA, 255–268. https://doi.org/10.1145/2661136.2661157

[15] Alex Groce and Martin Erwig. 2012. Finding Common Ground: Choose, Assert, and Assume. In *International Workshop on Dynamic Analysis*. 12–17.

[16] Alex Groce, Josselin Feist, Gustavo Grieco, and Michael Colburn. 2020. What are the Actual Flaws in Important Smart Contracts (and How Can We Find Them)?. In *International Conference on Financial Cryptography and Data Security*. 634–653.

[17] Alex Groce, Josie Holmes, Darko Marinov, August Shi, and Lingming Zhang. 2018. An Extensible, Regular-expression-based Tool for Multi-language Mutant Generation. In *Proceedings of the 40th International Conference on Software Engineering: Companion Proceedings* (Gothenburg, Sweden) *(ICSE '18)*. ACM, New York, NY, USA, 25–28. https://doi.org/10.1145/3183440.3183485

[18] Alex Groce, Gerard Holzmann, and Rajeev Joshi. 2007. Randomized Differential Testing as a Prelude to Formal Verification. In *29th International Conference on Software Engineering (ICSE'07)*. 621–631. https://doi.org/10.1109/ICSE.2007.68

[19] Alex Groce, Chaoqiang Zhang, Mohammad Amin Alipour, Eric Eide, Yang Chen, and John Regehr. 2013. Help, help, I'm being suppressed! The significance of suppressors in software testing. In *2013 IEEE 24th International Symposium on Software Reliability Engineering (ISSRE)*. IEEE, 390–399.

[20] Alex Groce, Chaoqiang Zhang, Eric Eide, Yang Chen, and John Regehr. 2012. Swarm Testing. In *International Symposium on Software Testing and Analysis*. 78–88.

[21] Josie Holmes, Iftekhar Ahmed, Caius Brindescu, Rahul Gopinath, He Zhang, and Alex Groce. 2020. Using Relative Lines of Code to Guide Automated Test Generation for Python. *ACM Trans. Softw. Eng. Methodol.* 29, 4, Article 28 (Sept. 2020), 38 pages. https://doi.org/10.1145/3408896

[22] Gerard J Holzmann. 2006. The power of 10: Rules for developing safety-critical code. *Computer* 39, 6 (2006), 95–99.

[23] George Klees, Andrew Ruef, Benji Cooper, Shiyi Wei, and Michael Hicks. 2018. Evaluating Fuzz Testing. In *Proceedings of the 2018 ACM SIGSAC Conference on Computer and Communications Security* (Toronto, Canada) *(CCS '18)*. Association for Computing Machinery, New York, NY, USA, 2123–2138. https://doi.org/10.1145/3243734.3243804

[24] Vu Le, Mehrdad Afshari, and Zhendong Su. 2014. Compiler validation via equivalence modulo inputs. *ACM SIGPLAN Notices* 49, 6 (2014), 216–226.

[25] William McKeeman. 1998. Differential testing for software. *Digital Technical Journal of Digital Equipment Corporation* 10(1) (1998), 100–107.

[26] Satoshi Nakamoto. 2008. Bitcoin: A peer-to-peer electronic cash system. https://bitcoin.org/bitcoin.pdf.

[27] Mike Papadakis, Marinos Kintis, Jie Zhang, Yue Jia, Yves Le Traon, and Mark Harman. 2019. Mutation testing advances: an analysis and survey. In *Advances in Computers*. Vol. 112. Elsevier, 275–378.

[28] Jingyu Zhou, Meng Xu, Alexander Shraer, Bala Namasivayam, Alex Miller, Evan Tschannen, Steve Atherton, Andrew J Beamon, Rusty Sears, John Leach, et al. 2021. FoundationDB: A Distributed Unbundled Transactional Key Value Store. In *ACM SIGMOD*.