

Looking for Lacunae in Bitcoin Core's Fuzzing Efforts

Alex Groce

Northern Arizona University
United States

ABSTRACT

Bitcoin is one of the most prominent distributed software systems in the world, and a key part of a potentially revolutionary new form of financial tool, cryptocurrency. At heart, Bitcoin exists as a set of nodes running an implementation of the Bitcoin protocol. This paper describes an effort to investigate and enhance the effectiveness of the Bitcoin Core implementation fuzzing effort. The effort initially began as a query about how to escape *saturation* in the fuzzing effort, but developed into a more general exploration once it was determined that saturation was largely illusory, a byproduct of the (then) fuzzing configuration. This paper reports the process and outcomes of the two-week focused effort that emerged from that initial contact between Chaincode labs and academic researchers. That effort found no smoking guns indicating major test/fuzz weaknesses. However, it produced a large number of additional fuzz corpus entries to add to the Bitcoin QA assets, clarified some long-standing problems in OSSFuzz triage, increased the set of documented fuzzers used in Bitcoin Core testing, and ran the first mutation analysis of Bitcoin Core code, revealing opportunities for further improvement. We contrast the Bitcoin Core transaction verification testing with similar tests for the most popular Ethereum and dogecoin implementations. This paper provides an overview of the challenges involved in improving testing infrastructure, processes, and documentation for a highly visible open source target system, from both the state-of-the-art research perspective and the practical engineering perspective.

CCS CONCEPTS

• **Software and its engineering** → **Dynamic analysis**; **Software testing and debugging**.

KEYWORDS

fuzzing, saturation, test diversity, mutation analysis

ACM Reference Format:

Alex Groce. 2021. Looking for Lacunae in Bitcoin Core's Fuzzing Efforts. In *Proceedings of the 2021 ACM SIGPLAN International Symposium on New Ideas, New Paradigms, and Reflections on Programming and Software (Onward! '21)*, October 20–22, 2021, Chicago, IL, USA. ACM, New York, NY, USA, 4 pages. <https://doi.org/10.1145/3486607.3486772>

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than ACM must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

Onward! '21, October 20–22, 2021, Chicago, IL, USA

© 2021 Association for Computing Machinery.

ACM ISBN 978-1-4503-9110-8/21/10...\$15.00

<https://doi.org/10.1145/3486607.3486772>

1 INTRODUCTION

Bitcoin [11] is the most popular cryptocurrency, and one of the most visible (if controversial) “new” systems based on software to rise to prominence in the last decade. As we write, while volatile, Bitcoin consistently has a market cap of over half a trillion dollars since January of 2021. Due to its distributed, decentralized nature, Bitcoin in some sense is the sum of the operations of the code executed by many independent Bitcoin nodes, especially nodes that mine cryptocurrency. Bitcoin Core (<https://github.com/Bitcoin/Bitcoin>) is by far the most popular implementation, and serves as a reference for all other implementations. To a significant degree, the code of Bitcoin Core is Bitcoin. The main Bitcoin Core repo on GitHub has over 57,000 stars, and has been forked more than 30,000 times.

Because of its fame and the high monetary value of Bitcoins, the Bitcoin protocol and its implementations are a high-value target for hackers (or even nation states interested in controlling cryptocurrency developments). Therefore, testing the code is of paramount importance, including extensive functional tests and aggressive *fuzzing*. This paper describes a focused effort to identify weaknesses in, and improve, the fuzzing of Bitcoin Core.

2 INITIAL CONTACT AND THE PROBLEM OF SATURATION

Chaincode labs (<https://chaincode.com/>) is a private R&D center based in Manhattan that exists solely to support and develop Bitcoin. In March of 2021, Adam Jonas, the head of special projects at Chaincode, contacted the first author to discuss determining a strategy to improve the fuzzing of Bitcoin Core. In particular, at the time, it seemed that the fuzzing was “stuck”: neither code coverage nor found bugs were increasing with additional fuzzing time. After some discussion, an 80 hour effort was determined as a reasonable scope for an external, research-oriented, look at the fuzzing effort. Before that effort, conducted over the summer of 2021, began, the problem of saturation resolved itself. Nonetheless, the issue that drove the initial desire for a researcher investigation is well worth examining. Moreover, understanding why Bitcoin Core fuzzing was, temporarily but not fundamentally, saturated, may be useful to help other fuzzing campaigns avoid the same false saturation problem.

Saturation, as defined in the blog post (<https://blog.regehr.org/archives/1796>) that brought Chaincode labs to the first author, is when “We apply a fuzzer to some non-trivial system... initially it finds a lot of bugs... [but] the number of new bugs found by the fuzzer drops off, eventually approaching zero.” That is, at first a fuzzer applied to a system will tend to continuously, sometimes impressively, increase both coverage and discovery of previously-unknown bugs. But, at some point, these bugs are known (and often fixed) and the fuzzer stops producing new bugs. Code and behavioral coverage seems to be *saturated*.

The underlying reason for saturation is that any fuzzer (or other test generator) explores a space of generated tests according to

some, perhaps very complex, probability distribution. Some bugs lie in the high-probability portion of this space, and other bugs like in very low probability (or in some cases, zero probability) parts of the space. Unsurprisingly, eventually the high probability space is well-explored, and the remaining bugs are found only very infrequently, if ever. The underlying empirical facts are cruel, as noted by Böhme and Falk: “[W]ith twice the machines, we can find *all known* bugs in half the time. Yet, finding linearly *more* bugs in the same time requires exponentially more machines” [2].

Chaincode labs saw that code coverage and bugs were not increasing in their fuzzer runs, and wanted to break out of the saturation trap. The blog post suggested several methods for doing just that, and this paper explores some of them. Note that even though the fuzzing was not saturated, the methods for escaping saturation are also useful things to try in any fuzzing effort where finding as many bugs as possible is actually important.

2.1 We Hold the World Ransom For... ONE MILLION FUZZER ITERATIONS

One problem for inexperienced fuzzer users is that it is not always clear just how long a fuzzer run is needed. Anyone used to popular random testing tools, such as QuickCheck [4], may expect a “reasonable” budget to be on the order of hundreds of tests or at most a few minutes [9]. The fuzz testing research community has settled, to some extent, on 24 hours as a basis for evaluation of fuzzers [10], but even this may be considered a low-budget run in a serious campaign! The Solidity compiler fuzzing effort discussed in the saturation blog post (<https://blog.trailofbits.com/2020/06/05/breaking-the-solidity-compiler-with-a-fuzzer/>, <https://blog.trailofbits.com/2021/03/23/a-year-in-the-life-of-a-compiler-fuzzing-campaign/>) found many bugs only after running a specialized version of AFL for over a month.

One thing that quickly emerged from discussions with Chaincode before the primary 80 hour effort was the limited extent of the fuzzer runs performed in early April. The fuzzing includes a large number of targets, each with its own fuzz harness and executable. At the time, the basic strategy was to run libFuzzer on each of these for 100,000 iterations. Because some targets are very fast and a few, such as full message processing, are slow, this meant in practice fuzzing most targets for only 30-90 seconds, and even the slowest targets for only a little over an hour. The total time for over 100 targets was not negligible, but expecting such short runs for each target, after an initial exploration of the easy part of the probability space, to gain coverage or bugs very often, was simply unrealistic. In particular, for complex, critical targets such as transaction verification and end-to-end message processing, 100,000 iterations was as inadequate as Dr. Evil’s famous ransom demand for “one million dollars” in the first Austen Powers movie. Neither the Chaincode team nor Dr. Evil was being unreasonable; they were both simply unaware of the “inflation” of needed resources required for serious fuzzing or blackmail of the world.

The first suggestion for escaping coverage, therefore was very simple: run the fuzzer longer! The Chaincode team immediately tried increasing their configuration to 5 million iterations (multiplying the number of executions, and runtime, by a factor of 50. Based on initial success with a few targets, this was done for all

targets, and eventually became the new default. To a large extent, saturation was no longer a problem. This exploration of simply increasing the fuzzing budget came early, about 15 days after the initial contact. The Chaincode team also added new seeds by, as advised, running more fuzzers, including AFL and Honggfuzz, in the same time frame.

By May 20th, Bitcoin Core was also in OSS-Fuzz (it was not at the time of the first discussions, due to reporting requirements, but negotiations settled this problem): <https://github.com/google/oss-fuzz/tree/master/projects/bitcoin-core>. From then on, Bitcoin Core has essentially been continuously fuzzed, and OSS-Fuzz quickly produced new crashes to investigate, and continues to do so: <https://bugs.chromium.org/p/oss-fuzz/issues/list?q=bitcoin>.

3 ADDING FUZZER DIVERSITY: USING ECLIPSE AND TRYING ENSEMBLE FUZZING

The most obvious solution when fuzzer A faces saturation on a target codebase is to bring in fuzzer B. In fact, the original blog post that attracted Chaincode’s attention mentions this in the very definition of saturation: “Subsequently, a different fuzzer, applied to the same system, finds a lot of bugs.” Of course, this is the easy case, and there can also be saturation across the meta-fuzzer defined as a composition of all fuzzers applied! Still, “throw in another fuzzer” is almost always a good idea when performing serious fuzzing campaigns. If that fuzzer adds some elements that are not present in any of the previously applied fuzzers, so much the better.

3.1 Fuzzing with Eclipser

Eclipser [3] is a fuzzer that combines afl-like¹ coverage-driven mutation-based fuzzing with a scalable form of grey-box concolic testing. AFL++ [5] and libFuzzer <https://llvm.org/docs/LibFuzzer.html>, the two fuzzers most aggressively applied to the Bitcoin Core code, do not perform any kind of symbolic or concolic testing. Eclipser is fairly easy to apply to new systems, since it uses QEMU to run on an ordinary binary, so the first order of business was to see if it could find new bugs or at least cover code AFL++ and libFuzzer were not able to reach.

Either directly (from runs using Eclipser itself) or indirectly (performing long AFL or libFuzzer runs seeded with Eclipser-generated tests), Eclipser produced thousands of new corpus seed files that were accepted as PRs to merge into the Bitcoin Core QA assets repo <https://github.com/bitcoin-core/qa-assets>. The first author is now the third-largest contributor to the QA assets repo, in fact. These files added hundreds of new coverage edges to the basis for OSS-Fuzz and other testing of Bitcoin Core.

On the other hand, these files added *no* new covered statements; the additions were all edges, not exploration of completely uncovered code, the most promising form of new coverage [1, 7]. Nor did these tests expose any new bugs. Of course, in the long run testing in OSS-Fuzz or on the QA team’s servers may use these tests to help cover new code or find new bugs, but the immediate impact is more modest.

¹In fact, the latest version of Eclipser simply uses afl for non concolic fuzzing.

One of the outcomes of the 80 hour effort was full documentation of how to run Eclipser v1.x on the Bitcoin Core code: <https://github.com/bitcoin/bitcoin/blob/master/doc/fuzzing.md#fuzzing-bitcoin-core-using-eclipser-v1x>.

4 SIDE ISSUES: TRIAGE, SPACE USAGE, AND FRUSTRATING BUGS

5 TRYING SWARM FUZZING

DeepState [6] provides strong support for swarm fuzzing. However, while there are similarities to the bitcoin core developed infrastructure for fuzz targets and the DeepState API, re-writing the fuzz targets to use DeepState is too large an effort for the likely payoff at this time, even though DeepState would also provide additional fuzzer diversity.

6 MUTATION ANALYSIS

We used the universal mutator (<https://github.com/agroce/universalmutator>) [8] to mutate the code.

6.1 Comparison with Other Cryptocurrency Codebases

7 HOW GOOD IS THE BITCOIN CORE TESTING AND FUZZING?

8 CONCLUSIONS

REFERENCES

- [1] Iftexhar Ahmed, Rahul Gopinath, Caius Brindescu, Alex Groce, and Carlos Jensen. 2016. Can Testedness Be Effectively Measured?. In *Proceedings of the 2016 24th ACM SIGSOFT International Symposium on Foundations of Software Engineering* (Seattle, WA, USA) (FSE 2016). Association for Computing Machinery, New York, NY, USA, 547–558. <https://doi.org/10.1145/2950290.2950324>
- [2] Marcel Böhme and Brandon Falk. 2020. Fuzzing: On the Exponential Cost of Vulnerability Discovery. In *Proceedings of the 28th ACM Joint Meeting on European Software Engineering Conference and Symposium on the Foundations of Software Engineering* (Virtual Event, USA) (ESEC/FSE 2020). Association for Computing Machinery, New York, NY, USA, 713–724. <https://doi.org/10.1145/3368089.3409729>
- [3] Jaeseung Choi, Joonun Jang, Choongwoo Han, and Sang Kil Cha. 2019. Grey-box Concolic Testing on Binary Code. In *Proceedings of the International Conference on Software Engineering*. 736–747.
- [4] Koen Claessen and John Hughes. 2000. QuickCheck: a lightweight tool for random testing of Haskell programs. In *International Conference on Functional Programming (ICFP)*. 268–279.
- [5] Andrea Fioraldi, Dominik Maier, Heiko Eißfeldt, and Marc Heuse. 2020. AFL++ : Combining Incremental Steps of Fuzzing Research. In *14th USENIX Workshop on Offensive Technologies (WOOT 20)*. USENIX Association. <https://www.usenix.org/conference/woot20/presentation/fioraldi>
- [6] Peter Goodman and Alex Groce. 2018. DeepState: Symbolic unit testing for C and C++. In *NDSS Workshop on Binary Analysis Research*.
- [7] Rahul Gopinath, Carlos Jensen, and Alex Groce. 2014. Code Coverage for Suite Evaluation by Developers. In *Proceedings of the 36th International Conference on Software Engineering* (Hyderabad, India) (ICSE 2014). Association for Computing Machinery, New York, NY, USA, 72–82. <https://doi.org/10.1145/2568225.2568278>
- [8] Alex Groce, Josie Holmes, Darko Marinov, August Shi, and Lingming Zhang. 2018. An Extensible, Regular-expression-based Tool for Multi-language Mutant Generation. In *Proceedings of the 40th International Conference on Software Engineering: Companion Proceedings* (Gothenburg, Sweden) (ICSE '18). ACM, New York, NY, USA, 25–28. <https://doi.org/10.1145/3183440.3183485>
- [9] Josie Holmes, Iftexhar Ahmed, Caius Brindescu, Rahul Gopinath, He Zhang, and Alex Groce. 2020. Using Relative Lines of Code to Guide Automated Test Generation for Python. *ACM Trans. Softw. Eng. Methodol.* 29, 4, Article 28 (Sept. 2020), 38 pages. <https://doi.org/10.1145/3408896>
- [10] George Klees, Andrew Ruef, Benji Cooper, Shiyi Wei, and Michael Hicks. 2018. Evaluating Fuzz Testing. In *Proceedings of the 2018 ACM SIGSAC Conference on Computer and Communications Security* (Toronto, Canada) (CCS '18). Association for Computing Machinery, New York, NY, USA, 2123–2138. <https://doi.org/10.1145/3243734.3243804>
- [11] Satoshi Nakamoto. 2008. Bitcoin: A peer-to-peer electronic cash system. <https://bitcoin.org/bitcoin.pdf>.




Figure 1: The Gaping Maw of Cthulhu